

Module 6: Process Synchronization

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Critical Regions
- Monitors
- Synchronization in Solaris 2
- Atomic Transactions

Background

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Shared-memory solution to bounded-buffer problem (Chapter 4) allows at most $n - 1$ items in buffer at the same time. A solution, where all N buffers are used is not simple.
 - Suppose that we modify the producer-consumer code by adding a variable *counter*, initialized to 0 and incremented each time a new item is added to the buffer

Bounded-Buffer

- Shared data **type** *item* = ... ;
 var *buffer* **array** [0..*n*-1] **of** *item*;
 in, out. 0..*n*-1;
 counter. 0..*n*;
 in, out, counter := 0;
- Producer process
 repeat
 ...
 produce an item in *nextp*
 ...
 while *counter* = *n* **do** no-op;
 buffer [*in*] := *nextp*;
 in := *in* + 1 **mod** *n*;
 counter := *counter* + 1;
 until false;

Bounded-Buffer (Cont.)

- Consumer process

```
repeat
    while counter = 0 do no-op;
    nextc := buffer [out];
    out := out + 1 mod n;
    counter := counter - 1;
    ...
    consume the item in nextc
    ...
until false;
```

- The statements:
 - *counter* := *counter* + 1;
 - *counter* := *counter* - 1;must be executed *atomically*.

The Critical-Section Problem

- n processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.
- Structure of process P_i

repeat

entry section

critical section

exit section

reminder section

until *false*;

Solution to Critical-Section Problem

1. **Mutual Exclusion.** If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.
2. **Progress.** If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely.
3. **Bounded Waiting.** A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.
 - Assume that each process executes at a nonzero speed
 - No assumption concerning relative speed of the n processes.

Initial Attempts to Solve Problem

- Only 2 processes, P_0 and P_1
- General structure of process P_i (other process P_j)

repeat

entry section

critical section

exit section

reminder section

until *false*;

- Processes may share some common variables to synchronize their actions.

Algorithm 1

- Shared variables:
 - **var** *turn*: (0..1);
initially *turn* = 0
 - *turn* = *i* \Rightarrow P_i can enter its critical section
- Process P_i

repeat

while *turn* \neq *i* **do** *no-op*;
critical section

turn := *j*;

reminder section

until *false*;

- Satisfies mutual exclusion, but not progress

Algorithm 2

- Shared variables
 - **var** *flag*: **array** [0..1] **of** *boolean*;
initially *flag* [0] = *flag* [1] = *false*.
 - *flag* [*i*] = *true* \Rightarrow P_i ready to enter its critical section
- Process P_i

repeat

flag[*i*] := *true*;
while *flag*[*j*] **do** *no-op*;

critical section

flag [*i*] := *false*;

remainder section

until *false*;

- Satisfies mutual exclusion, but not progress requirement.

Algorithm 3

- Combined shared variables of algorithms 1 and 2.
- Process P_i

repeat

flag [i] := true;

turn := j;

while (*flag [j]* and *turn = j*) **do** *no-op*;

critical section

flag [i] := false;

remainder section

until *false*;

- Meets all three requirements; solves the critical-section problem for two processes.

Bakery Algorithm

Critical section for n processes

- Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- If processes P_i and P_j receive the same number, if $i < j$, then P_i is served first; else P_j is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...

Bakery Algorithm (Cont.)

- Notation \leq lexicographical order (ticket #, process id #)
 - $(a,b) < c,d$ if $a < c$ or if $a = c$ and $b < d$
 - $\max(a_0, \dots, a_{n-1})$ is a number, k , such that $k \geq a_i$ for $i = 0, \dots, n - 1$
- Shared data

var *choosing*: **array** $[0..n - 1]$ **of** *boolean*;

number. **array** $[0..n - 1]$ **of** *integer*,

Data structures are initialized to *false* and 0 respectively

Bakery Algorithm (Cont.)

repeat

```
    choosing[i] := true;  
    number[i] := max(number[0], number[1], ..., number [n − 1])+1;  
    choosing[i] := false;  
    for j := 0 to n − 1  
        do begin  
            while choosing[j] do no-op;  
            while number[j] ≠ 0  
                and (number[j], j) < (number[i], i) do no-op;  
        end;
```

critical section

```
number[i] := 0;
```

remainder section

until *false*;

Synchronization Hardware

- Test and modify the content of a word atomically.

```
function Test-and-Set (var target: boolean): boolean;  
    begin  
        Test-and-Set := target;  
        target := true;  
    end;
```

Mutual Exclusion with Test-and-Set

- Shared data: **var** *lock*: *boolean* (*initially false*)
- Process P_i

repeat

while *Test-and-Set* (*lock*) **do** *no-op*;

critical section

lock := *false*;

remainder section

until *false*;

Semaphore

- Synchronization tool that does not require busy waiting.
- Semaphore S – integer variable
- can only be accessed via two indivisible (atomic) operations

wait (S): **while** $S \leq 0$ **do** *no-op*;
 $S := S - 1$;

signal (S): $S := S + 1$;

Example: Critical Section of n Processes

- Shared variables
 - **var** *mutex* : *semaphore*
 - initially *mutex* = 1
- Process P_i

repeat

wait(mutex);

critical section

signal(mutex);

remainder section

until *false*;

Semaphore Implementation

- Define a semaphore as a record

type semaphore = record

value: integer

L: list of process;

end;

- Assume two simple operations:
 - block suspends the process that invokes it.
 - wakeup(*P*) resumes the execution of a blocked process *P*.

Implementation (Cont.)

- Semaphore operations now defined as

wait(S): $S.value := S.value - 1;$

if $S.value < 0$

then begin

 add this process to $S.L$;

block;

end;

signal(S): $S.value := S.value + 1;$

if $S.value \leq 0$

then begin

 remove a process P from $S.L$;

wakeup(P);

end;

Semaphore as General Synchronization Tool

- Execute B in P_j only after A executed in P_i
- Use semaphore $flag$ initialized to 0
- Code:

P_i	P_j
\vdots	\vdots
A	$wait(flag)$
$signal(flag)$	B

Deadlock and Starvation

- Deadlock – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let S and Q be two semaphores initialized to 1

P_0	P_1
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
\vdots	\vdots
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q)</i>	<i>signal(S);</i>

- Starvation – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Two Types of Semaphores

- *Counting* semaphore – integer value can range over an unrestricted domain.
- *Binary* semaphore – integer value can range only between 0 and 1; can be simpler to implement.
- Can implement a counting semaphore S as a binary semaphore.

Implementing S as a Binary Semaphore

- Data structures:

```
var S1: binary-semaphore;  
    S2: binary-semaphore;  
    S3: binary-semaphore;  
    C: integer;
```

- Initialization:

```
S1 = S3 = 1  
S2 = 0  
C = initial value of semaphore S
```

Implementing S (Cont.)

- *wait* operation

```
wait(S3);  
wait(S1);  
C := C - 1;  
if C < 0  
then begin  
    signal(S1);  
    wait(S2);  
end  
else signal(S1);  
    signal(S3);
```

- *signal* operation

```
wait(S1);  
C := C + 1;  
if C ≤ 0 then signal(S2);  
    signal(S1);
```


Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded-Buffer Problem

- Shared data

```
type item = ...  
var buffer = ...  
    full, empty, mutex: semaphore;  
    nextp, nextc: item;  
    full := 0; empty := n; mutex := 1;
```

Bounded-Buffer Problem (Cont.)

- Producer process

repeat

...

produce an item in *nextp*

...

wait(empty);

wait(mutex);

...

signal(mutex);

signal(full);

until *false*;

Bounded-Buffer Problem (Cont.)

- Consumer process

repeat

wait(full)

wait(mutex);

...

remove an item from *buffer* to *nextc*

...

signal(mutex);

signal(empty);

...

consume the item in *nextc*

...

until *false*;

Readers-Writers Problem

- Shared data

```
var mutex, wrt. semaphore (=1);  
    readcount : integer (=0);
```

- Writer process


```
wait(wrt);  
    ...  
    writing is performed  
    ...  
signal(wrt);
```

Readers-Writers Problem (Cont.)

- Reader process

```
wait(mutex);  
    readcount := readcount + 1;  
    if readcount = 1 then wait(wrt);  
signal(mutex);  
    ...  
    reading is performed  
    ...  
wait(mutex);  
    readcount := readcount - 1;  
    if readcount = 0 then signal(wrt);  
signal(mutex);
```

Dining-Philosophers Problem



- Shared data

var chopstick: array [0..4] of semaphore;
(=1 initially)

Dining-Philosophers Problem (Cont.)

- Philosopher i :

```
repeat
    wait(chopstick[i])
    wait(chopstick[i+1 mod 5])
    ...
    eat
    ...
    signal(chopstick[i]);
    signal(chopstick[i+1 mod 5]);
    ...
    think
    ...
until false;
```


Critical Regions

- High-level synchronization construct
- A shared variable v of type T , is declared as:

var v : shared T

- Variable v accessed only inside statement

region v when B do S

where B is a Boolean expression.

While statement S is being executed, no other process can access variable v .

Critical Regions (Cont.)

- Regions referring to the same shared variable exclude each other in time.
- When a process tries to execute the region statement, the Boolean expression B is evaluated. If B is true, statement S is executed. If it is false, the process is delayed until B becomes true and no other process is in the region associated with v .

Example – Bounded Buffer

- Shared variables:

var *buffer*: **shared record**

pool: **array** [0..*n*−1] **of** *item*;
count,*in*,*out*: **integer**

end;

- Producer process inserts *nextp* into the shared buffer

region *buffer* **when** *count* < *n*
do begin

pool[*in*] := *nextp*;
in := *in* + 1 **mod** *n*;
count := *count* + 1;

end;

Bounded Buffer Example (Cont.)

- Consumer process removes an item from the shared buffer and puts it in nextc

```
region buffer when count > 0  
  do begin  
    nextc := pool[out];  
    out := out+1 mod n;  
    count := count - 1;  
  end;
```

Implementation: region x when B do S

- Associate with the shared variable x , the following variables:
var mutex, first-delay, second-delay: semaphore;
first-count, second-count: integer,
- Mutually exclusive access to the critical section is provided by *mutex*.
- If a process cannot enter the critical section because the Boolean expression **B** is false, it initially waits on the *first-delay semaphore*; moved to the *second-delay semaphore* before it is allowed to reevaluate B .

Implementation (Cont.)

- Keep track of the number of processes waiting on *first-delay* and *second-delay*, with *first-count* and *second-count* respectively.
- The algorithm assumes a FIFO ordering in the queuing of processes for a semaphore.
- For an arbitrary queuing discipline, a more complicated implementation is required.

```

wait(mutex);
while not B
  do begin first-count := first-count + 1;
    if second-count > 0
      then signal(second-delay)
      else signal(mutex);
    wait(first-delay);
    first-count := first-count - 1;
    if first-count > 0 then signal(first-delay)
      else signal(second-delay);
    wait(second-delay);
    second-count := second-count - 1;
  end;

S;
if first-count > 0
  then signal(first-delay);
  else if second-count > 0
    then signal(second-delay);
    else signal(mutex);

```

Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.

```
type monitor-name = monitor
    variable declarations
    procedure entry P1 : (...);
        begin ... end;
    procedure entry P2(...);
        begin ... end;
        ⋮
    procedure entry Pn (...);
        begin...end;
    begin
        initialization code
    end
```


Monitors (Cont.)

- To allow a process to wait within the monitor, a *condition* variable must be declared, as

var *x, y: condition*

- Condition variable can only be used with the operations *wait* and *signal*.
 - The operation


x.wait;

means that the process invoking this operation is
suspended until another process invokes


x.signal;

- The *x.signal* operation resumes exactly one suspended process. If no process is suspended, then the signal operation has no effect.

Schematic view of a monitor



Monitor with condition variables



Dining Philosophers Example

```
type dining-philosophers = monitor
  var state : array [0..4] of :(thinking, hungry, eating);
  var self : array [0..4] of condition;
  procedure entry pickup (i: 0..4);
  begin
    state[i] := hungry,
    test (i);
    if state[i] ≠ eating then self[i], wait,
  end;

  procedure entry putdown (i: 0..4);
  begin
    state[i] := thinking;
    test (i+4 mod 5);
    test (i+1 mod 5);
  end;
```

Dining Philosophers (Cont.)

```
procedure test(k: 0..4);
begin
    if state[k+4 mod 5] ≠ eating
       and state[k] = hungry
       and state[k+1 mod 5] ≠ eating
    then begin
        state[k] := eating;
        self[k].signal;
    end;

end;

begin
    for i := 0 to 4
        do state[i] := thinking;
    end.

end.
```

Monitor Implementation Using Semaphores

- Variables

```
var mutex: semaphore (init = 1)
    next: semaphore (init = 0)
    next-count: integer (init = 0)
```

- Each external procedure F will be replaced by
wait(mutex);

...

body of F ;

...

```
if next-count > 0
    then signal(next)
    else signal(mutex);
```

- Mutual exclusion within a monitor is ensured.

Monitor Implementation (Cont.)

- For each condition variable x , we have:
 var x -sem: semaphore (init = 0)
 x -count: integer (init = 0)
- The operation x .wait can be implemented as:

```
 $x$ -count :=  $x$ -count + 1;  
if next-count > 0  
    then signal(next)  
    else signal(mutex);  
wait( $x$ -sem);  
 $x$ -count :=  $x$ -count - 1;
```

Monitor Implementation (Cont.)

- The operation `x.signal` can be implemented as:

```
if x-count > 0
  then begin
    next-count := next-count + 1;
    signal(x-sem);
    wait(next);
    next-count := next-count - 1;
  end;
```


Monitor Implementation (Cont.)

- *Conditional-wait* construct: *x.wait(c);*
 - *c* – integer expression evaluated when the wait operation is executed.
 - value of *c* (*priority number*) stored with the name of the process that is suspended.
 - when *x.signal* is executed, process with smallest associated priority number is resumed next.
- Check two conditions to establish correctness of system:
 - User processes must always make their calls on the monitor in a correct sequence.
 - Must ensure that an uncooperative process does not ignore the mutual-exclusion gateway provided by the monitor, and try to access the shared resource directly, without using the access protocols.

Solaris 2 Operating System

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing.
- Uses adaptive mutexes for efficiency when protecting data from short code segments.
- Uses condition variables and readers-writers locks when longer sections of code need access to data.