

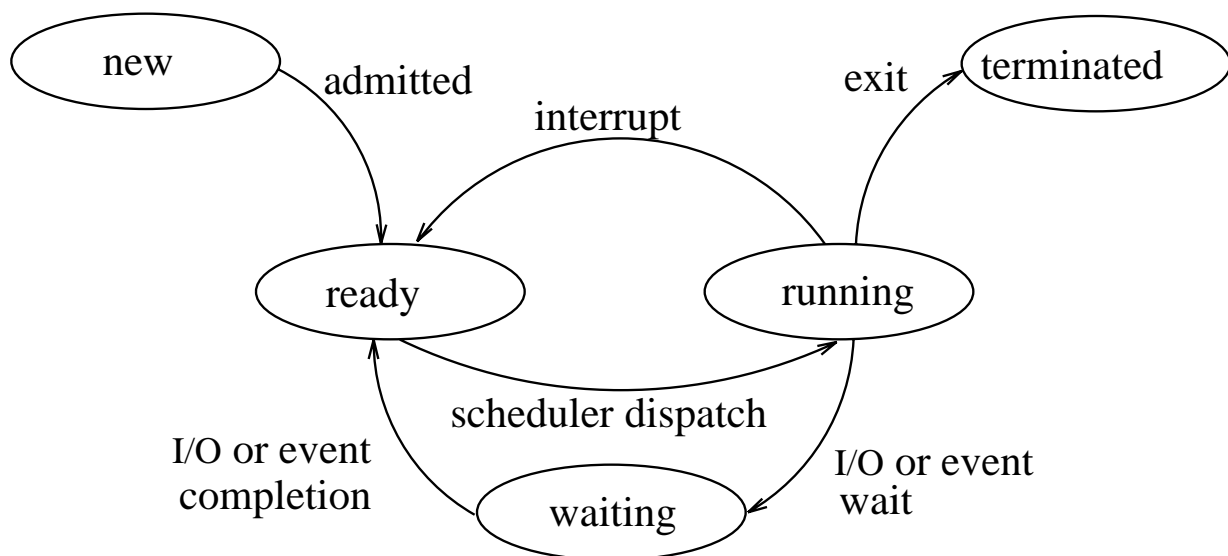
CHAPTER 4: PROCESSES

- Process Concept
- Process Scheduling
- Operation on Processes
- Cooperating Processes
- Threads
- Interprocess Communication

Process Concept


- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably.
- Process – a program in execution; process execution must progress in a sequential fashion.
- A process includes:
 - *program counter*
 - *stack*
 - *data section*

- As a process executes, it changes *state*.
 - **New:** The process is being created.
 - **Running:** Instructions are being executed.
 - **Waiting:** The process is waiting for some event to occur.
 - **Ready:** The process is waiting to be assigned to a processor.
 - **Terminated:** The process has finished execution.
- Diagram of process state:



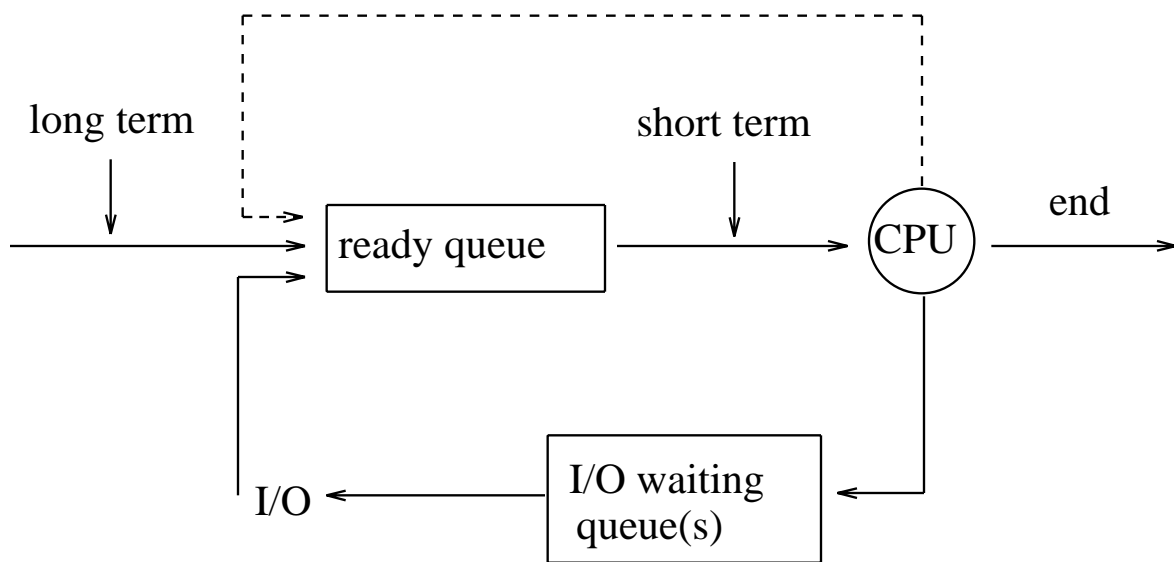
- Process Control Block (PCB) – Information associated with each process.
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information

- Process scheduling queues
 - *job queue* – set of all processes in the system.
 - *ready queue* – set of all processes residing in main memory, ready and waiting to execute.
 - *device queues* – set of processes waiting for a particular I/O device.
- Process migration between the various queues.



- Schedulers

- *Long-term scheduler (job scheduler)* – selects which processes should be brought into the ready queue.
- *Short-term scheduler (CPU scheduler)* – selects which process should be executed next and allocates CPU.



- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow).
- The long-term scheduler controls the *degree of multiprogramming*.
- Processes can be described as either:
 - *I/O-bound process* – spends more time doing I/O than computations; many short CPU bursts.
 - *CPU-bound process* – spends more time doing computations; few very long CPU bursts.

Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
- Time dependent on hardware support.

Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes.
- Resource sharing
 - Parent and children share all resources.
 - Children share subset of parent's resources.
 - Parent and child share no resources.
- Execution
 - Parent and children execute concurrently.
 - Parent waits until children terminate.
- Address space
 - Child duplicate of parent.
 - Child has a program loaded into it.
- UNIX examples
 - **fork** system call creates new process.
 - **execve** system call used after a **fork** to replace the process' memory space with a new program.

Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**).
 - Output data from child to parent (via **fork**).
 - Process' resources are deallocated by operating system.
- Parent may terminate execution of children processes (**abort**).
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting.
 - Operating system does not allow child to continue if its parent terminates.
 - *Cascading termination.*

Cooperating Processes

- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process.
- Advantages of process cooperation:
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Producer-Consumer Problem

- Paradigm for cooperating processes; *producer* process produces information that is consumed by a *consumer* process.
 - *unbounded-buffer* places no practical limit on the size of the buffer.
 - *bounded-buffer* assumes that there is a fixed buffer size.
- Shared-memory solution:
 - Shared data

```
var n;  
type item = ... ;  
var buffer: array [0..n-1] of item;  
in, out: 0..n-1;  
in := 0;  
out := 0;
```

- Producer process

repeat

...

produce an item in *nextp*

...

while $in+1 \bmod n = out$ **do** *no-op*;

buffer[in] := nextp;

in := in+1 mod n;

until *false*;

- Consumer process

repeat

while $in = out$ **do** *no-op*;

nextc := buffer[out];

out := out+1 mod n;

...

consume the item in *nextc*

...

until *false*;

- Solution is correct, but can only fill up $n-1$ buffer.

Threads

- A *thread* (or *lightweight process*) is a basic unit of CPU utilization; it consists of:
 - program counter
 - register set
 - stack space
- A thread shares with its peer threads its:
 - code section
 - data section
 - operating-system resourcescollectively known as a *task*.
- A traditional or *heavyweight* process is equal to a task with one thread.

- In a task containing multiple threads, while one server thread is blocked and waiting, a second thread in the same task could run.
 - Cooperation of multiple threads in same job confers higher throughput and improved performance.
 - Applications that require sharing a common buffer (producer–consumer problem) benefit from thread utilization.
- Threads provide a mechanism that allows sequential processes to make blocking system calls while also achieving parallelism.
- Kernel-supported threads (Mach and OS/2).
- User-level threads; supported above the kernel, via a set of library calls at the user level (Project Andrew from CMU).
- Hybrid approach implements both user-level and kernel-supported threads (Solaris 2).

Solaris 2 – version of UNIX with support for threads at the kernel and user levels, symmetric multiprocessing, and real-time scheduling.

- LWP – intermediate level between user-level threads and kernel-level threads.
- Resource needs of thread types:
 - Kernel thread – small data structure and a stack; thread switching does not require changing memory access information, and therefore is relatively fast.
 - LWP – PCB with register data, accounting information, and memory information; switching between LWPs is relatively slow.
 - User-level thread – needs only a stack and a program counter. Switching is fast since kernel is not involved. Kernel only sees the LWPs in the process that support user-level threads.

Interprocess Communication (IPC) – provides a mechanism to allow processes to communicate and to synchronize their actions.

- Message system – processes communicate with each other without resorting to shared variables.
- IPC facility provides two operations:
 - **send**(*message*) – messages can be of either fixed or variable size.
 - **receive**(*message*)
- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Communication link
 - physical implementation (e.g., shared memory, hardware bus)
 - logical implementation (e.g., logical properties)

Implementation questions:

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bidirectional?

Direct Communication

- Processes must name each other explicitly:
 - **send**(P , *message*) – send a message to process P
 - **receive**(Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically.
 - A link is associated with exactly one pair of communicating processes.
 - Between each pair there exists exactly one link.
 - The link may be unidirectional, but is usually bidirectional.

Indirect Communication

- Messages are directed and received from *mailboxes* (also referred to as *ports*).
 - Each mailbox has a unique *id*.
 - Processes can communicate only if they share a mailbox.
- Properties of communication link
 - Link established only if the two processes share a mailbox in common.
 - A link may be associated with many processes.
 - Each pair of processes may share several communication links.
 - Link may be unidirectional or bidirectional.
- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox

Indirect Communication (Continued)

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A.
 - P_1 sends; P_2 and P_3 receive.
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Buffering – queue of messages attached to the link; implemented in one of three ways.

- Zero capacity – 0 messages
Sender must wait for receiver (*rendezvous*).
- Bounded capacity – finite length of n messages
Sender must wait if link full.
- Unbounded capacity – infinite length
Sender never waits.

Exception Conditions – error recovery

- Process terminates
- Lost messages
- Scrambled Messages

