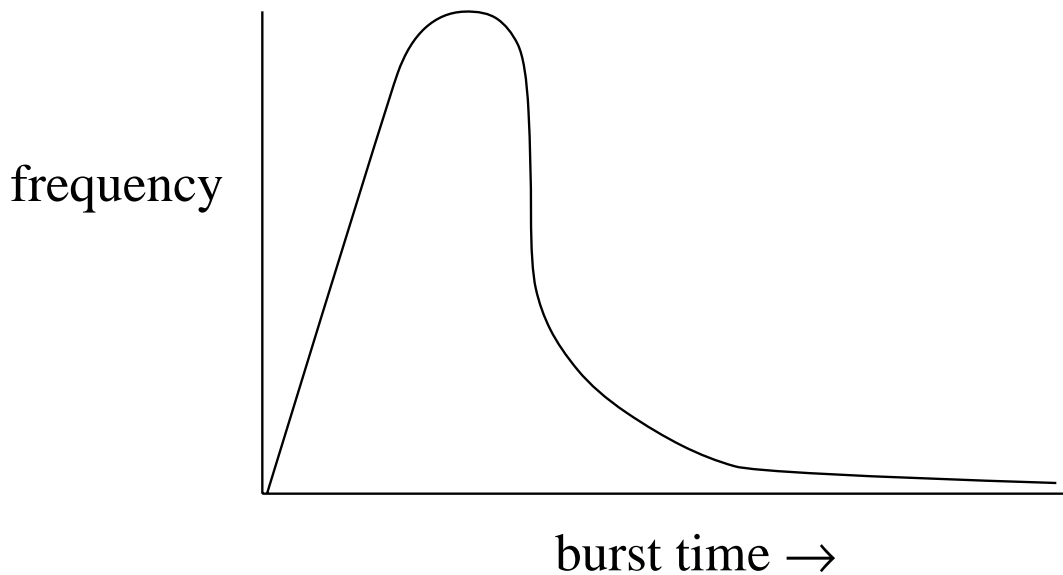


## CHAPTER 5: CPU SCHEDULING

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Real-Time Scheduling
- Algorithm Evaluation

## Basic Concepts

- Maximum CPU utilization obtained with multi-programming.
- CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait.
- CPU burst distribution



- *Short-term scheduler* –selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
  1. switches from running to waiting state.
  2. switches from running to ready state.
  3. switches from waiting to ready.
  4. terminates.
- Scheduling under 1 and 4 is *nonpreemptive*.
- All other scheduling is *preemptive*.

# Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- *Dispatch latency* – time it takes for the dispatcher to stop one process and start another running.

## Scheduling Criteria

- CPU utilization – keep the CPU as busy as possible
- Throughput – # of processes that complete their execution per time unit
- Turnaround time – amount of time to execute a particular process
- Waiting time – amount of time a process has been waiting in the ready queue
- Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

- Optimization
  - Max CPU utilization
  - Max throughput
  - Min turnaround time
  - Min waiting time
  - Min response time

## First-Come, First-Served (FCFS) Scheduling

- Example:

<u>Process</u>	<u>Burst time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:

$P_1, P_2, P_3$ .

The Gantt chart for the schedule is:

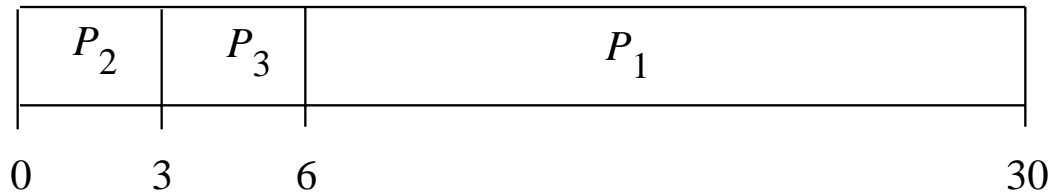


- Waiting time for:  
 $P_1 = 0$   
 $P_2 = 24$   
 $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

- Suppose that the processes arrive in the order:

$$P_2, P_3, P_1.$$

The Gantt chart for the schedule is:



- Waiting time for:
 
$$P_1 = 6$$

$$P_2 = 0$$

$$P_3 = 3$$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case.
- *Convoy effect*: short process behind long process



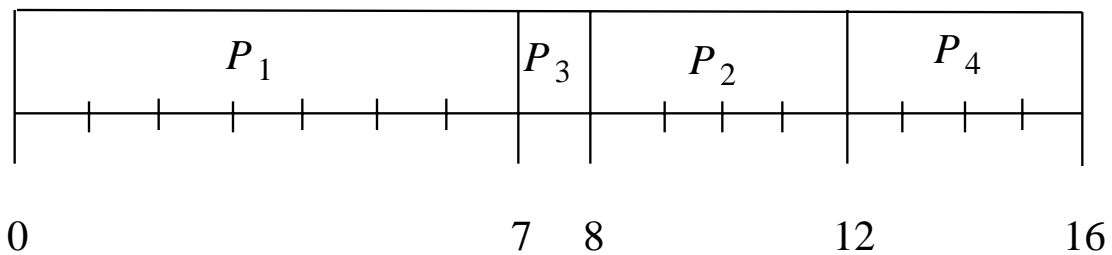
## Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- Two schemes:
  - a) nonpreemptive – once CPU given to the process it cannot be preempted until it completes its CPU burst.
  - b) preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF).
- SJF is optimal – gives minimum average waiting time for a given set of processes.

## Example of SJF

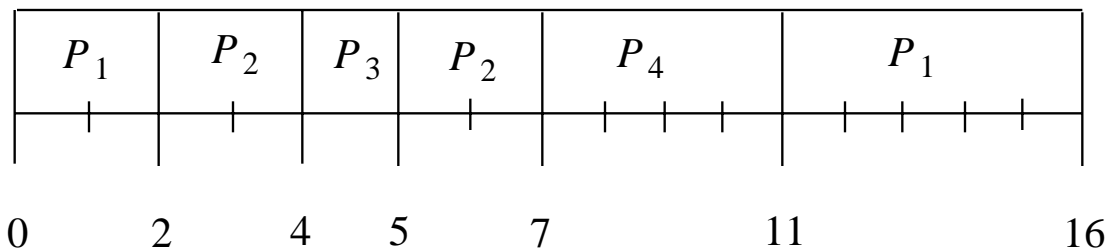
- | <u>Process</u> | <u>Arrival time</u> | <u>CPU time</u> |
|----------------|---------------------|-----------------|
| $P_1$          | 0                   | 7               |
| $P_2$          | 2                   | 4               |
| $P_3$          | 4                   | 1               |
| $P_4$          | 5                   | 4               |

- SJF (non-preemptive)



$$\text{Average waiting time} = (0 + 6 + 3 + 7)/4 = 4$$

- SRTF (preemptive)



$$\text{Average waiting time} = (9 + 1 + 0 + 2)/4 = 3$$

How do we know the length of the next CPU burst?

- Can only estimate the length.
- Can be done by using the length of previous CPU bursts, using exponential averaging.

1.  $T_n$  = actual length of  $n^{\text{th}}$  CPU burst

2.  $\psi_n$  = predicted value of  $n^{\text{th}}$  CPU burst

3.  $0 \leq W \leq 1$

4. Define:

$$\psi_{n+1} = W * T_n + (1 - W) \psi_n$$

## Examples:

- $W = 0$

$$\Psi_{n+1} = \Psi_n$$

Recent history does not count.

- $W = 1$

$$\Psi_{n+1} = T_n$$

Only the actual last CPU burst counts.

- If we expand the formula, we get:

$$\begin{aligned}\Psi_{n+1} = & W * T_n + (1 - W) * W * T_{n-1} + \\ & (1 - W)^2 * W * T_{n-2} + \dots + (1 - W)^q \\ & * W * T_{n-q}\end{aligned}$$

So if  $W = 1/2 \Rightarrow$  each successive term has less and less weight.

## Priority Scheduling

- A priority number (integer) is associated with each process.
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority).
  - a) preemptive
  - b) nonpreemptive
- SJN is a priority scheduling where priority is the predicted next CPU burst time.
- Problem  $\equiv$  Starvation – low priority processes may never execute.

Solution  $\equiv$  Aging – as time progresses increase the priority of the process.

## Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10–100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Performance

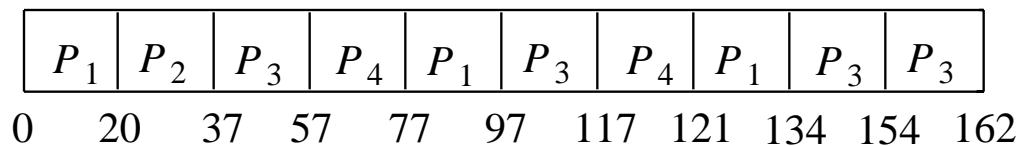
$q$  large  $\Rightarrow$  FIFO

$q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high.

## Example of RR with time quantum = 20

- | <u>Process</u> | <u>CPU times</u> |
|----------------|------------------|
| $P_1$          | 53               |
| $P_2$          | 17               |
| $P_3$          | 68               |
| $P_4$          | 24               |

- The Gantt chart is:



- Typically, higher average turnaround than SRT, but better *response*.

## Multilevel Queue

- Ready queue is partitioned into separate queues.

Example: foreground (interactive)  
background (batch)

- Each queue has its own scheduling algorithm.

Example: foreground – RR  
background – FCFS

- Scheduling must be done between the queues.

- Fixed priority scheduling

Example: serve all from foreground then from background. Possibility of starvation.

- Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes.

Example:

80% to foreground in RR

20% to background in FCFS



## Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way.
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithm for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

## Example of multilevel feedback queue

- Three queues:
  - $Q_0$  – time quantum 8 milliseconds
  - $Q_1$  – time quantum 16 milliseconds
  - $Q_2$  – FCFS

- Scheduling

A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ . At  $Q_1$ , job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .

- Multiple-Processor Scheduling
  - CPU scheduling more complex when multiple CPUs are available.
  - *Homogeneous* processors within a multiprocessor.
  - *Load sharing*
  - *Asymmetric multiprocessing* – only one processor accesses the system data structures, alleviating the need for data sharing.
- Real-Time Scheduling
  - *Hard real-time* systems – required to complete a critical task within a guaranteed amount of time.
  - *Soft real-time* computing – requires that critical processes receive priority over less fortunate ones.

## Algorithm Evaluation

- *Deterministic modeling* – takes a particular predetermined workload and defines the performance of each algorithm for that workload.
- Queueing models
- Implementation

