

CHAPTER 9: VIRTUAL MEMORY

- Background
- Demand Paging
- Performance of Demand Paging
- Page Replacement
- Page-Replacement Algorithms
- Allocation of Frames
- Thrashing
- Other Considerations
- Demand Segmentation

Background

- Virtual memory – separation of user logical memory from physical memory.
 - Only *part* of the program needs to be in memory for execution.
 - Logical address space can therefore be much larger than physical address space.
 - Need to allow pages to be *swapped* in and out.
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation

Demand Paging

- Bring a page into memory only when it is needed.
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory

Valid–Invalid bit

- With each page table entry a valid–invalid bit is associated ($1 \Rightarrow$ in-memory, $0 \Rightarrow$ not-in-memory)
- Initially valid–invalid bit is set to 0 on all entries.
- Example of a page table snapshot.


| frame # | valid-invalid bit |
|---------|-------------------|
| | 1 |
| | 1 |
| | 1 |
| | 1 |
| | 0 |
| ⋮ | |
| | 0 |
| | 0 |

page table

- During address translation, if valid–invalid bit in page table entry is $0 \Rightarrow$ page fault.

Page Fault

1. If there is ever a reference to a page, first reference will trap to OS \Rightarrow *page fault*.
2. OS looks at another table to decide:
 - a) Invalid reference \Rightarrow abort.
 - b) Just not in memory.
3. Get empty frame.
4. Swap page into frame.
5. Reset tables, validation bit = 1.
6. Restart instruction:
 - block move



- auto increment/decrement location

What happens if there is no free frame?

- Page replacement – find some page in memory, but not really in use, swap it out.

⇒ algorithm

⇒ performance – want an algorithm which will result in minimum number of page faults.

- Same page may be brought into memory several times.

Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$, no page faults
 - if $p = 1$, every reference is a fault

- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p \text{ (page fault overhead} \\ & + [\text{swap page out}] \\ & + \text{swap page in} \\ & + \text{restart overhead}) \end{aligned}$$

- Example:

- memory access time = 1 microsecond
- 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out.
- Swap Page Time = 10 msec = 10,000 msec
- $\text{EAT} = (1 - p) \times 1 + p (15000)$
 $= 1 + 15000P \quad (\text{in msec})$

Page Replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement.
- Use *modify (dirty) bit* to reduce overhead of page transfers – only modified pages are written to disk.
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory.

Page-Replacement Algorithms

- Want lowest *page-fault rate*.
- Evaluate algorithm by running it on a particular string of memory references (*reference string*) and computing the number of page faults on that string.
- In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.

First-In-First-Out (FIFO) Algorithm

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

- 3 frames (3 pages can be in memory at a time per process)

| | | | | |
|---|---|---|---|---------------|
| 1 | 1 | 4 | 5 | 9 page faults |
| 2 | 2 | 1 | 3 | |
| 3 | 3 | 2 | 4 | |

- 4 frames

| | | | | |
|---|---|---|---|----------------|
| 1 | 1 | 5 | 4 | 10 page faults |
| 2 | 2 | 1 | 5 | |
| 3 | 3 | 2 | | |
| 4 | 4 | 3 | | |

FIFO Replacement – Belady’s Anomaly

more frames \nRightarrow less page faults

Optimal Algorithm

- Replace the page that will not be used for the longest period of time.
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

| | |
|---|---|
| 1 | 4 |
| 2 | |
| 3 | |
| 4 | 5 |

6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs.

Least Recently Used (LRU) Algorithm

| | | |
|---|---|---|
| 1 | | 5 |
| 2 | | |
| 3 | 5 | 4 |
| 4 | 3 | |

- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter.
 - When a page needs to be changed, look at the counters to determine which are to change
- Stack implementation – keep a stack of page numbers in a double link form:
 - Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
 - No search for replacement

LRU Approximation Algorithms

- Reference bit
 - With each page associate a bit, initially = 0.
 - When page is referenced bit set to 1.
 - Replace the one which is 0 (if one exists). We do not know the order, however.
- Second chance
 - Need reference bit.
 - Clock replacement.
 - If page to be replaced (in clock order) has reference bit = 1, then:
 - a) set reference bit 0.
 - b) leave page in memory.
 - c) replace next page (in clock order), subject to same rules.

- Counting Algorithms – keep a counter of the number of references that have been made to each page.
 - LFU Algorithm: replaces page with smallest count.
 - MFU Algorithm: based on the argument that the page with the smallest count was probably just brought in and has yet to be used.
- Page-Buffering Algorithm – desired page is read into a free frame from the *pool* before the victim is written out.

Allocation of Frames

- Each process needs minimum number of pages.

Example: IBM 370 – 6 pages to handle
SS MOVE instruction:

- a) Instruction is 6 bytes, might span 2 pages.
 - b) 2 pages to handle **from**.
 - c) 2 pages to handle **to**.
- Two major allocation schemes:
 - fixed allocation
 - priority allocation

- Fixed allocation

- Equal allocation

If 100 frames and 5 processes, give each 20 pages.

- Proportional allocation

Allocate according to the size of process.

- s_i = size of process p_i
 - $S = \sum s_i$
 - m = total number of frames
 - a_i = allocation for $p_i = \frac{s_i}{S} \times m$

Example : $m = 64$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$


- Priority allocation
 - Use a proportional allocation scheme using priorities rather than size.
 - If process P_i generates a page fault,
 - select for replacement one of its frames.
 - select for replacement a frame from a process with lower priority number.

Global versus local allocation

- Global replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- Local replacement – each process selects from only its own set of allocated frames.

Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high:
 - ⇒ low CPU utilization.
 - ⇒ operating system thinks that it needs to increase the degree of multiprogramming.
 - ⇒ another process added to the system.
- Thrashing \equiv a process is busy swapping pages in and out.



- Why does paging work?

Locality model

- Process migrates from one locality to another.
- Localities may overlap.

- Why does thrashing occur?

Σ size of locality > total memory size

Working-Set Model

- $\Delta \equiv$ working-set window \equiv a fixed number of page references

Example: 10,000 instruction

- WSS_i – working set of process $P_i =$
total number of pages referenced in the most recent Δ (varies in time)

If Δ too small will not encompass entire locality.

If Δ too large will encompass several localities.

If $\Delta = \infty \Rightarrow$ will encompass entire program.

- $D = \Sigma WSS_i \equiv$ total demand frames
- If $D > m \Rightarrow$ thrashing.
- Policy if $D > m$, then suspend one of the processes.

How do you keep track of the working set?

- Approximate with:

interval timer + a reference bit

- Example:


$$\Delta = 10,000$$

- Timer interrupts after every 5000 time units.
- Keep in memory 2 bits for each page.
- Whenever a timer interrupts copy and sets the values of all reference bits to 0.
- If one of the bits in memory = 1 \Rightarrow page in working set.

Not completely accurate (why?)

Improve = 10 bits and interrupt every 1000 time units

Page-Fault Frequency Scheme



- Establish “acceptable” page-fault rate.
 - If actual rate too low, process loses frame.
 - If actual rate too high, process gains frame.

Other Considerations

1. Prepaging

2. Page size selection

- fragmentation
- table size
- I/O overhead
- locality

3. Program structure

- Array $A[1024,1024]$ of integer
- Each row is stored in one page
- One frame
- Program 1 **for** $j := 1$ to 1024 **do**
 for $i := 1$ to 1024 **do**
 $A[i, j] := 0;$

1024 × 1024 page faults

- Program 2 **for** $i := 1$ to 1024 **do**
 for $j := 1$ to 1024 **do**
 $A[i, j] := 0;$

1024 page faults

4. I/O interlock and addressing

Demand Segmentation – used when insufficient hardware to implement demand paging.

- OS/2 allocates memory in segments, which it keeps track of through *segment descriptors*.
- Segment descriptor contains a valid bit to indicate whether the segment is currently in memory.
 - If segment is in main memory, access continues,
 - If not in memory, segment fault.

